

LLM Cost Tracker by PHANTSOM

Business & Non-Technical Overview

Author: Soham Dahivalkar | PHANTSOM

Date: May 2026

The Problem

Every company using AI (ChatGPT, Claude, Gemini) in their software products pays money **per API call**. These costs add up fast:

Scenario	Monthly Cost
A chatbot answering 1,000 customer queries/day using GPT-4	\$5,400/month
The same chatbot using GPT-4o-mini instead	\$54/month
Savings from one model switch	\$5,346/month

The problem? Developers writing the code have **zero visibility** into what each AI call costs. They pick a model, write the code, and the finance team gets a shock at the end of the month.

There is no tool that shows the cost **while writing code** — before it hits production.

Until now.

The Solution

LLM Cost Tracker by PHANTSOM is a tool that plugs into the developer's code editor (VS Code — used by 15+ million developers worldwide) and shows the **estimated cost of every AI API call** in real-time, right next to the code.

What the Developer Sees:

BEFORE (without our tool):

model="gpt-4" ← Developer has no idea this costs \$0.09 per call

AFTER (with our tool):

model="gpt-4" ← 💀 \$0.09/call · ~\$270/mo [GPT-4] – EXPENSIVE!

💡 Consider gpt-4o-mini – 97% quality at 3% of the cost

The developer **immediately knows**:

- This call costs \$0.09
- At 100 calls/day, that's \$270/month
- There's a cheaper option that's almost as good

Key Features

1. Real-Time Cost Visibility

Every line of code that calls an AI model gets an automatic cost label — **color-coded by expense level**:

Color	Meaning	Action
🟢 Green	Very cheap	No worries
🟡 Yellow	Moderate	Monitor it
🟠 Orange	Expensive	Review the choice
🔴 Red	Very expensive	Find alternatives
💀 Skull	Dangerously expensive	Immediate attention needed

2. Detailed Cost Breakdown (on hover)

When a developer hovers their mouse over any AI model name, they see:

- Cost per single call
- Daily / Weekly / Monthly / Yearly projections
- Full pricing details from the AI provider
- Context window and model specifications

3. Automatic Cheaper Alternative Suggestions

If a developer is using an expensive model, the tool **automatically suggests a cheaper one**:

- "You're using GPT-4 (\$30/1M tokens). Consider GPT-4o-mini (\$0.15/1M tokens) — **200x cheaper**"
- "You're using Claude 3 Opus (\$15/1M). Consider Claude 3 Haiku (\$0.25/1M) — **60x cheaper**"

4. Monthly Budget Alerts

Teams can set a monthly AI budget. If estimated costs exceed the budget, the tool shows a **warning notification — before the code goes to production.**

5. Multi-Currency Support

Costs can be displayed in:

- us USD (US Dollar)
- eu EUR (Euro)
- gb GBP (British Pound)
- in INR (Indian Rupee)
- jp JPY (Japanese Yen)

6. Comprehensive AI Provider Coverage

Supports **40+ models** from **8 major AI providers**:

Provider	Example Models
OpenAI	GPT-4o, GPT-4, GPT-3.5, o1, o3, o4-mini
Anthropic	Claude 3 Opus, Claude 3.5 Sonnet, Claude 4
Google	Gemini 2.5 Pro, Gemini 2.0 Flash
Mistral	Mistral Large, Codestral
DeepSeek	DeepSeek V3, DeepSeek R1
Cohere	Command R, Command R+
Groq	LLaMA 3.1, Mixtral
AWS	Amazon Titan Text

Who Is This For?

User	How They Benefit
Software Developers	See costs before deploying code. Make informed model choices.
Engineering Managers	Ensure teams stay within AI budgets. Prevent surprise bills.
CTOs / Tech Leads	Get visibility into AI infrastructure costs at the code level.
Startup Founders	Control AI spending from Day 1. Every dollar matters.
Freelance AI Developers	Accurately estimate costs for client projects.
Students / Learners	Understand the financial impact of different AI models.

Business Impact

Cost Savings

Without This Tool	With This Tool
Developer picks GPT-4 "because it's the best"	Developer sees GPT-4 costs \$270/mo and GPT-4o-mini costs \$2.70/mo
Nobody notices until the monthly bill arrives	Cost is visible at the moment of writing code
Finance team asks "why did AI costs triple?"	Budget alerts trigger before code reaches production
Result: Uncontrolled AI spending	Result: Optimized AI spending from Day 1

Real-World Savings Example

A typical AI-powered SaaS product might have:

- 5 AI features
- Each making 500 API calls/day
- Using GPT-4o (\$2.50 input / \$10.00 output per 1M tokens)

Without optimization: ~\$450/month

After switching 3 features to GPT-4o-mini (suggested by our tool): ~\$85/month

Annual savings: \$4,380

For larger companies with dozens of AI features, savings can reach **\$50,000-\$500,000+ per year.**

Competitive Landscape

Feature	LLM Cost Tracker by PHANTSOM	Manual Tracking (Spreadsheets)	Cloud Billing Dashboards
Shows cost while writing code	✔ Yes	✗ No	✗ No
Real-time (updates as you type)	✔ Yes	✗ No	✗ No (delayed by hours/days)
Suggests cheaper alternatives	✔ Yes	✗ No	✗ No
Supports 40+ models, 8 providers	✔ Yes	⚠ Manual effort	⚠ Only their own models
Budget alerts before deployment	✔ Yes	✗ No	⚠ After the fact
Free to use	✔ Yes	✔ Yes	✔ Yes
No internet/API keys needed	✔ Yes	N/A	✗ Requires account
Setup time	Instant (1 click install)	Hours	Minutes to Hours

Market Opportunity

The AI API Market is Exploding

- **\$15+ billion** spent annually on AI API calls globally (2025)
- **15+ million** developers use VS Code (our target platform)

- **Every company** building AI products needs cost control
- **No dominant tool** exists in this space yet

Target Market Size

Segment	Size	Our Reach
VS Code Users	15M+ developers	Available to all
AI/ML Developers	3M+ globally	Primary audience
Companies using AI APIs	500,000+	Enterprise potential
AI Startups	50,000+	Early adopter audience

Product Specifications

Specification	Details
Platform	VS Code (Windows, Mac, Linux)
Installation	One-click from VS Code Marketplace
Size	25 KB (extremely lightweight)
Performance Impact	Near zero — doesn't slow down the editor
Internet Required	No — works fully offline
Languages Supported	Python, TypeScript, JavaScript, React
AI Models Covered	40+ models from 8 providers
Price	Free (MIT License)
Privacy	No data collected. No telemetry. Everything runs locally.

Distribution & Availability

Channel	Status
VS Code Marketplace	✔ Published — searchable by 15M+ developers
Direct Install	✔ Available via .vsix file
Open Source	✔ MIT License — free for personal and commercial use

Marketplace Link:

`https://marketplace.visualstudio.com/items?itemName=sohamdahivalkar.llm-cost-estimator`

Future Roadmap

Phase	Feature	Business Impact
Phase 2	Team dashboard — shared budget tracking	Enterprise sales
Phase 3	CI/CD integration — block deployments over budget	DevOps pipeline
Phase 4	Historical cost tracking & analytics	Cost optimization insights
Phase 5	Support for more languages (Java, Go, Rust, C#)	Wider developer reach
Phase 6	Live pricing updates (fetch real-time prices)	Always accurate costs
Phase 7	Custom model pricing (self-hosted models)	Enterprise customization

About the Creator

Soham Dahivalkar — Generative AI Engineer

Achievement	Details
Current Role	Generative AI Engineer at Alembic Pharmaceuticals
Published Book	"Generative AI: High Stakes Cyber Security" (Amazon Kindle)
Research Paper	"AI in Security: ML Approach for Vulnerability Management" (ResearchGate)
Open Source Library	ai-bridge-kit (Published on PyPI)
VS Code Extension	LLM Cost Tracker by PHANTSOM (VS Code Marketplace)
Education	B.E. Computer Engineering — 9.52 SGPA
Recognition	State-Level Top 3 Finalist — RTMSU Startup Competition
Certifications	Microsoft + IIT Bombay (AI/ML), IBM (AI Fundamentals), and 4 more

Brand: PHANTSOM — *"The Invisible Force Behind Intelligent Systems"*

Contact

Channel	Details
Email	sohamdahivalkar4@gmail.com
LinkedIn	linkedin.com/in/soham-dahivalkar-82415426a
GitHub	github.com/phantsom

Summary

LLM Cost Tracker by PHANTSOM solves a simple but critical problem: **developers have no idea how much their AI API calls cost while writing code.** Our tool puts a price tag on every AI call — instantly, visually, and intelligently — with cheaper alternatives suggested automatically.

One sentence pitch:

"We help developers see AI costs before they become bills."

Built with purpose by PHANTSOM — Soham Dahivalkar

"The Invisible Force Behind Intelligent Systems"